

# ANZE XIE

(608)-236-3958 | [a1xie@ucsd.edu](mailto:a1xie@ucsd.edu) | [AnzeXie.github.io](https://AnzeXie.github.io)

## EDUCATION

<b>University of California-San Diego</b> M.S. in Computer Science & Engineering	<i>Sept. 2022 - Jun. 2024</i> GPA: 3.9/4.0
<b>University of Wisconsin-Madison</b> B.S. in Computer Science & Statistics & Applied Mathematics	<i>Sept. 2018 - Dec. 2021</i> GPA: 3.99/4.00

## PUBLICATIONS

Lanxiang Hu\*, Qiyu Li\*, **Anze Xie\***, Nan Jiang, Ion Stoica, Haojian Jin, and Hao Zhang. "[GameArena: Evaluating LLM Reasoning through Live Computer Games](#)." (ICLR 2025). Co-first authored.

Dacheng Li\*, Rulin Shao\*, **Anze Xie**, Eric P. Xing, Joseph E. Gonzalez, Ion Stoica, Xuezhe Ma, and Hao Zhang. "[DISTFLASHATTN: Distributed Memory-efficient Attention for Long-context LLMs Training](#)." (COLM 2024).

Dacheng Li\*, Rulin Shao\*, **Anze Xie**, Ying Sheng, Lianmin Zheng, Joseph E. Gonzalez, Ion Stoica, Xuezhe Ma, and Hao Zhang. "[How long can opensource llms truly promise on context length](#)." (Instruction workshop @ NeurIPS 2023).

**Anze Xie**, Anders Carlsson, Jason Mohoney, Roger Waleffe, Shanan Peters, Theodoros Rekatsinas, and Shivaram Venkataraman. "[Demo of marius: a system for large-scale graph embeddings](#)." *Proceedings of the VLDB Endowment* 14, no. 12 (2021): 2759-2762.

## RESEARCH EXPERIENCE

<b>Research Assistant</b> <i>Advisor: <a href="#">Prof. Hao Zhang</a></i> Area of Research: Large language models and corresponding training systems	UCSD, USA <i>Mar. 2023 - Present</i>
<b>Research Assistant</b> <i>Advisor: <a href="#">Prof. Shivaram Venkataraman</a>, <a href="#">Prof. Theodoros Rekatsinas</a></i> Area of Research: Systems for graph learning	UW-Madison, USA <i>Feb. 2021 - May. 2022</i>
<b>Research Intern</b> <i>Advisor: <a href="#">Prof. Anne-Marie Kermarrec</a></i> Area of Research: Federated learning	EPFL, Switzerland <i>May. 2022 - Aug. 2022</i>

## OPEN-SOURCE CONTRIBUTIONS

Contributed significantly to [LongChat](#) and [LongEval](#), a repository supports training and evaluation of long-context LLMs

Developed the data preprocessing, postprocessing, and rule-based configuration optimizer modules for [Marius](#) and [MariusGNN](#), a unified system for large-scale graph-learning tasks

## TEACHING EXPERIENCE

<b>Teaching Assistant</b> Scalable Data Systems ML Systems	UCSD, USA <i>Jan. 2024 - Mar. 2024</i> <i>Mar. 2024 - Jun. 2024</i>
--	---

## AWARDS

UW-Madison Undergraduate Scholarship for Summer Study, 2020, 2021

UW-Madison Dean's List (7 semesters)

## SKILLS

**Languages:** Python, GO, C, C++, Java, SQL, Matlab, R, HTML/CSS,  $\text{\LaTeX}$

**Frameworks and Tools:** PyTorch, Tensorflow, Transformers, pytest, tox